

ConsistEval_Experiments

May 22, 2026

1 ConsistEval: Augmenting Summarisation Evaluation with NLI-Based Factual Consistency Verification

SIT330-770 Natural Language Processing — High Distinction Task 2.3

Author: Alexander Petkovski — Deakin University

This notebook implements the full experimental pipeline for the ConsistEval framework: 1. Load the SummEval benchmark dataset 2. Compute surface-level metrics (ROUGE-1, ROUGE-2, ROUGE-L, BERTScore) 3. Implement NLI-based consistency scoring using `roberta-large-mnli` 4. Compute the composite ConsistEval score 5. Calculate Spearman correlations with human expert consistency ratings 6. Generate results table and scatter plot 7. Perform error analysis

Runtime: Set your Colab runtime to **GPU** (Runtime → Change runtime type → T4 GPU)

1.1 1. Install Dependencies

```
[1]: !pip install -q rouge-score bert-score transformers torch nltk scipy pandas ↵  
    ↪matplotlib seaborn tqdm datasets
```

Preparing metadata (setup.py) ... done

61.1/61.1 kB

4.8 MB/s eta 0:00:00

Building wheel for rouge-score (setup.py) ... done

1.2 2. Imports and Setup

```
[2]: import json  
import urllib.request  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from scipy.stats import spearmanr  
from rouge_score import rouge_scorer  
from bert_score import score as bert_score_fn  
from transformers import AutoTokenizer, AutoModelForSequenceClassification  
import torch  
import nltk
```

```

from tqdm import tqdm
import warnings
warnings.filterwarnings('ignore')

nltk.download('punkt', quiet=True)
nltk.download('punkt_tab', quiet=True)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f'Using device: {device}')

```

Using device: cuda

1.3 3. Load the SummEval Benchmark

SummEval (Fabbri et al., 2021) contains 100 CNN/DailyMail source articles, each summarised by 16 different models (1,600 summary-article pairs). Each summary has expert human ratings for coherence, consistency, fluency, and relevance on a 1-5 Likert scale.

```

[3]: # Download SummEval data from the official repository
SUMMEVAL_URL = 'https://storage.googleapis.com/sfr-summarization-repo-research/
↳model_annotations.aligned.jsonl'

print('Downloading SummEval benchmark data...')
urllib.request.urlretrieve(SUMMEVAL_URL, 'summeval.jsonl')

# Parse the JSONL file
data = []
with open('summeval.jsonl', 'r') as f:
    for line in f:
        data.append(json.loads(line.strip()))

print(f'Loaded {len(data)} summary-article pairs')

# Inspect one entry to understand the structure
print(f'\nKeys in each entry: {list(data[0].keys())}')
print(f'Number of expert annotators: {len(data[0].get("expert_annotators",
↳[]))}')

```

Downloading SummEval benchmark data...

Loaded 1600 summary-article pairs

Keys in each entry: ['id', 'decoded', 'expert_annotators',
'turker_annotators', 'references', 'model_id', 'filepath']

Number of expert annotators: 3

1.3.1 3.1 Load CNN/DailyMail Source Articles

The SummEval JSONL file contains summaries, references, and human annotations, but **not** the original source articles. We load these separately from the HuggingFace `cnn_dailymail` dataset

and match them to SummEval entries by article ID.

```
[4]: from datasets import load_dataset

# Load CNN/DailyMail test set from HuggingFace
print('Loading CNN/DailyMail test set from HuggingFace (this may take a minute).
↳..')
cnn_dm = load_dataset('cnn_dailymail', '3.0.0', split='test')
print(f'Loaded {len(cnn_dm)} CNN/DailyMail test articles')

# Create a mapping from article ID (raw hash) to article text
id_to_article = {article['id']: article['article'] for article in cnn_dm}
print(f'Built ID-to-article mapping with {len(id_to_article)} entries')

# --- KEY FIX: SummEval IDs have a prefix like "dm-test-" or "cnn-test-" ---
# Example SummEval ID: 'dm-test-8764fb95bfad8ee849274873a92fb8d6b400eee2'
# Example CNN/DM ID: '8764fb95bfad8ee849274873a92fb8d6b400eee2'
# The filepath field confirms: 'cnndm/dailymail/stories/
↳8764fb95bfad8ee849274873a92fb8d6b400eee2.story'
# We extract the raw hash from the filepath for reliable matching.

def extract_hash_from_entry(entry):
    """Extract the raw article hash from SummEval entry."""
    # Primary: extract from filepath (most reliable)
    filepath = entry.get('filepath', '')
    if filepath:
        filename = filepath.split('/')[-1].replace('.story', '')
        return filename
    # Fallback: strip known prefixes from the id field
    entry_id = entry.get('id', '')
    for prefix in ['dm-test-', 'cnn-test-', 'dm-val-', 'cnn-val-', 'dm-train-',
↳'cnn-train-']:
        if entry_id.startswith(prefix):
            return entry_id[len(prefix):]
    return entry_id

# Test the matching
sample_hashes = [extract_hash_from_entry(entry) for entry in data[:5]]
sample_cnndm_ids = list(id_to_article.keys())[:5]
print(f'\nExtracted hashes from SummEval: {sample_hashes[:3]}')
print(f'Sample CNN/DM IDs: {sample_cnndm_ids[:3]}')

# Verify match rate
unique_hashes = set(extract_hash_from_entry(entry) for entry in data)
matched = unique_hashes.intersection(set(id_to_article.keys()))
print(f'\nUnique SummEval articles: {len(unique_hashes)}')
print(f'Matched to CNN/DM: {len(matched)}/{len(unique_hashes)}')
```

```

if len(matched) == 0:
    raise RuntimeError(
        f'FATAL: Could not match ANY SummEval articles to CNN/DM.\n'
        f'Sample SummEval hash: {sample_hashes[0]}\n'
        f'Sample CNN/DM ID: {sample_cnndm_ids[0]}\n'
        f'The ID formats may have changed. Manual inspection needed.'
    )
elif len(matched) < len(unique_hashes):
    print(f'WARNING: {len(unique_hashes) - len(matched)} articles could not be_
    ↪matched.')
else:
    print('All 100 articles matched successfully!')

```

Loading CNN/DailyMail test set from HuggingFace (this may take a minute)...

README.md: 0.00B [00:00, ?B/s]

Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.

WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.

3.0.0/train-00000-of-00003.parquet: 0%| | 0.00/257M [00:00<?, ?B/s]

3.0.0/train-00001-of-00003.parquet: 0%| | 0.00/257M [00:00<?, ?B/s]

3.0.0/train-00002-of-00003.parquet: 0%| | 0.00/259M [00:00<?, ?B/s]

3.0.0/validation-00000-of-00001.parquet: 0%| | 0.00/34.7M [00:00<?, ?
 ↪B/s]

3.0.0/test-00000-of-00001.parquet: 0%| | 0.00/30.0M [00:00<?, ?B/s]

Generating train split: 0%| | 0/287113 [00:00<?, ? examples/s]

Generating validation split: 0%| | 0/13368 [00:00<?, ? examples/s]

Generating test split: 0%| | 0/11490 [00:00<?, ? examples/s]

Loaded 11490 CNN/DailyMail test articles

Built ID-to-article mapping with 11490 entries

Extracted hashes from SummEval: ['8764fb95bfad8ee849274873a92fb8d6b400eee2',
 '8764fb95bfad8ee849274873a92fb8d6b400eee2',
 '8764fb95bfad8ee849274873a92fb8d6b400eee2']
 Sample CNN/DM IDs: ['f001ec5c4704938247d27a44948eebb37ae98d01',
 '230c522854991d053fe98a718b1defa077a8efef',
 '4495ba8f3a340d97a9df1476f8a35502bcce1f69']

Unique SummEval articles: 100

Matched to CNN/DM: 100/100

All 100 articles matched successfully!

```
[5]: # Extract source articles, summaries, references, and human consistency ratings
sources = []
summaries = []
references = []
human_consistency = []
human_coherence = []
human_fluency = []
human_relevance = []
model_ids = []
missing_sources = 0

for entry in data:
    # --- FIX: Match using extracted hash, not raw SummEval ID ---
    article_hash = extract_hash_from_entry(entry)
    src = id_to_article.get(article_hash, '')

    if not src:
        missing_sources += 1

    summ = entry.get('decoded', entry.get('summary', ''))
    ref = entry.get('references', [])
    if isinstance(ref, list):
        ref = ref[0] if ref else ''

    # Extract expert annotations (average across annotators)
    expert_anns = entry.get('expert_annotations', [])
    if expert_anns:
        cons_scores = [a.get('consistency', 0) for a in expert_anns]
        coh_scores = [a.get('coherence', 0) for a in expert_anns]
        flu_scores = [a.get('fluency', 0) for a in expert_anns]
        rel_scores = [a.get('relevance', 0) for a in expert_anns]
        avg_consistency = np.mean(cons_scores)
        avg_coherence = np.mean(coh_scores)
        avg_fluency = np.mean(flu_scores)
        avg_relevance = np.mean(rel_scores)
    else:
        avg_consistency = entry.get('consistency', 0)
        avg_coherence = entry.get('coherence', 0)
        avg_fluency = entry.get('fluency', 0)
        avg_relevance = entry.get('relevance', 0)

    sources.append(src)
    summaries.append(summ)
    references.append(ref)
```

```

human_consistency.append(avg_consistency)
human_coherence.append(avg_coherence)
human_fluency.append(avg_fluency)
human_relevance.append(avg_relevance)
model_ids.append(entry.get('model_id', 'unknown'))

human_consistency = np.array(human_consistency)
human_coherence = np.array(human_coherence)
human_fluency = np.array(human_fluency)
human_relevance = np.array(human_relevance)

print(f'Total samples: {len(sources)}')
print(f'Sources successfully loaded: {len(sources) - missing_sources}/
↳{len(sources)}')
if missing_sources > 0:
    print(f'WARNING: {missing_sources} entries have no source article!')
print(f'Sample source length: {len(sources[0])} characters')
print(f'Sample source preview: {sources[0][:200]}...')
print(f'\nConsistency rating distribution: mean={human_consistency.mean():.2f},
↳std={human_consistency.std():.2f}')
print(f'Rating range: [{human_consistency.min():.1f}, {human_consistency.max():.
↳1f}]')

# Sanity check: assert that sources are not empty
assert len(sources[0]) > 100, "ERROR: Source articles are still empty! Check
↳the ID matching above."
print('\nSanity check PASSED: Source articles loaded successfully.')

```

Total samples: 1600
 Sources successfully loaded: 1600/1600
 Sample source length: 1843 characters
 Sample source preview: Paul Merson has restarted his row with Andros Townsend after the Tottenham midfielder was brought on with only seven minutes remaining in his team's 0-0 draw with Burnley on Sunday. 'Just been watchin...

Consistency rating distribution: mean=4.66, std=0.92
 Rating range: [1.0, 5.0]

Sanity check PASSED: Source articles loaded successfully.

1.4 4. Compute Surface-Level Metrics

1.4.1 4.1 ROUGE Scores

ROUGE (Lin, 2004) computes n-gram overlap between the generated summary and the reference summary.

```
[6]: # Compute ROUGE-1, ROUGE-2, ROUGE-L F1 scores
scorer = rouge_scorer.RougeScorer(['rouge1', 'rouge2', 'rougeL'],
    ↪use_stemmer=True)

rouge1_scores = []
rouge2_scores = []
rougeL_scores = []

print('Computing ROUGE scores...')
for summ, ref in tqdm(zip(summaries, references), total=len(summaries)):
    scores = scorer.score(ref, summ)
    rouge1_scores.append(scores['rouge1'].fmeasure)
    rouge2_scores.append(scores['rouge2'].fmeasure)
    rougeL_scores.append(scores['rougeL'].fmeasure)

rouge1_scores = np.array(rouge1_scores)
rouge2_scores = np.array(rouge2_scores)
rougeL_scores = np.array(rougeL_scores)

print(f'ROUGE-1 mean: {rouge1_scores.mean():.4f}')
print(f'ROUGE-2 mean: {rouge2_scores.mean():.4f}')
print(f'ROUGE-L mean: {rougeL_scores.mean():.4f}')
```

Computing ROUGE scores...

100%| | 1600/1600 [00:01<00:00, 1324.82it/s]

ROUGE-1 mean: 0.4254

ROUGE-2 mean: 0.1980

ROUGE-L mean: 0.2869

1.4.2 4.2 BERTScore

BERTScore (Zhang et al., 2020) computes cosine similarity between contextualised token embeddings from a pretrained model.

```
[7]: # Compute BERTScore using roberta-large
print('Computing BERTScore (this may take a few minutes on GPU)...')
P, R, F1 = bert_score_fn(
    summaries,
    references,
    model_type='roberta-large',
    device=device,
    verbose=True,
    batch_size=32
)
```

```
bertscore_f1 = F1.numpy()
print(f'BERTScore F1 mean: {bertscore_f1.mean():.4f}')
```

Computing BERTScore (this may take a few minutes on GPU)...

```
config.json: 0%|          | 0.00/482 [00:00<?, ?B/s]
tokenizer_config.json: 0%|          | 0.00/25.0 [00:00<?, ?B/s]
vocab.json: 0.00B [00:00, ?B/s]
merges.txt: 0.00B [00:00, ?B/s]
tokenizer.json: 0.00B [00:00, ?B/s]
model.safetensors: 0%|          | 0.00/1.42G [00:00<?, ?B/s]
Loading weights: 0%|          | 0/389 [00:00<?, ?it/s]
```

RobertaModel LOAD REPORT from: roberta-large

Key	Status
lm_head.dense.bias	UNEXPECTED
lm_head.layer_norm.bias	UNEXPECTED
lm_head.layer_norm.weight	UNEXPECTED
lm_head.bias	UNEXPECTED
lm_head.dense.weight	UNEXPECTED
roberta.embeddings.position_ids	UNEXPECTED
pooler.dense.weight	MISSING
pooler.dense.bias	MISSING

Notes:

- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
- MISSING :those params were newly initialized because missing from the checkpoint. Consider training on your downstream task.

calculating scores...

computing bert embedding.

```
0%|          | 0/52 [00:00<?, ?it/s]
```

computing greedy matching.

```
0%|          | 0/50 [00:00<?, ?it/s]
```

done in 7.11 seconds, 225.07 sentences/sec

BERTScore F1 mean: 0.8644

1.5 5. Implement NLI-Based Consistency Scoring

This is the core of our proposed method. For each summary, we: 1. Segment the summary into sentences 2. Segment the source document into sentences 3. For each summary sentence, find

the maximum entailment probability across all source sentences 4. Average these scores to get the document-level NLI consistency score

We use `roberta-large-mnli` as our NLI model (Williams et al., 2018).

```
[8]: # Load the NLI model
print('Loading roberta-large-mnli for NLI inference...')
nli_model_name = 'roberta-large-mnli'
nli_tokenizer = AutoTokenizer.from_pretrained(nli_model_name)
nli_model = AutoModelForSequenceClassification.from_pretrained(nli_model_name).
    ↪to(device)
nli_model.eval()

# The model outputs logits for: [contradiction, neutral, entailment]
ENTAILMENT_IDX = 2

print('NLI model loaded successfully.')
```

Loading roberta-large-mnli for NLI inference...

config.json: 0%| | 0.00/688 [00:00<?, ?B/s]

tokenizer_config.json: 0%| | 0.00/25.0 [00:00<?, ?B/s]

vocab.json: 0.00B [00:00, ?B/s]

merges.txt: 0.00B [00:00, ?B/s]

tokenizer.json: 0.00B [00:00, ?B/s]

model.safetensors: 0%| | 0.00/1.43G [00:00<?, ?B/s]

Loading weights: 0%| | 0/393 [00:00<?, ?it/s]

RobertaForSequenceClassification LOAD REPORT from: roberta-large-mnli

Key	Status		
-----+-----+---+			
roberta.pooler.dense.weight	UNEXPECTED		
roberta.pooler.dense.bias	UNEXPECTED		

Notes:

- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.

NLI model loaded successfully.

```
[9]: def get_entailment_prob(premise, hypothesis, tokenizer, model, device,
    ↪max_length=512):
    """
    Compute the entailment probability  $P(\text{entailment} \mid \text{premise}, \text{hypothesis})$ 
    using the NLI model.

    Args:
```

```

        premise: The source sentence (what we check against)
        hypothesis: The summary sentence (what we are verifying)
        tokenizer: The NLI tokenizer
        model: The NLI model
        device: torch device
        max_length: Maximum token length for the input

Returns:
    float: Probability of entailment
    """
    inputs = tokenizer(
        premise, hypothesis,
        return_tensors='pt',
        truncation=True,
        max_length=max_length,
        padding=True
    ).to(device)

    with torch.no_grad():
        outputs = model(**inputs)
        probs = torch.softmax(outputs.logits, dim=-1)
        entailment_prob = probs[0, ENTAILMENT_IDX].item()

    return entailment_prob

def compute_nli_score(source, summary, tokenizer, model, device):
    """
    Compute the NLI-based consistency score for a source-summary pair.

    Following SummaC (Laban et al., 2022), we:
    1. Segment the source into sentences
    2. Segment the summary into sentences
    3. For each summary sentence, compute max entailment across source sentences
    4. Average across summary sentences

    Args:
        source: The source document text
        summary: The generated summary text
        tokenizer: The NLI tokenizer
        model: The NLI model
        device: torch device

    Returns:
        float: NLI consistency score in [0, 1]
    """
    # Segment into sentences

```

```

try:
    source_sents = nltk.sent_tokenize(source)
    summary_sents = nltk.sent_tokenize(summary)
except Exception:
    # Fallback: split on periods
    source_sents = [s.strip() for s in source.split('.') if s.strip()]
    summary_sents = [s.strip() for s in summary.split('.') if s.strip()]

if not summary_sents or not source_sents:
    return 0.0

# For each summary sentence, find max entailment across source sentences
sentence_scores = []
for summ_sent in summary_sents:
    if len(summ_sent.strip()) < 5: # Skip very short fragments
        continue

    max_entailment = 0.0
    for src_sent in source_sents:
        if len(src_sent.strip()) < 5:
            continue
        ent_prob = get_entailment_prob(
            src_sent, summ_sent, tokenizer, model, device
        )
        max_entailment = max(max_entailment, ent_prob)

    sentence_scores.append(max_entailment)

if not sentence_scores:
    return 0.0

# Document-level score: mean across summary sentences
return np.mean(sentence_scores)

```

```

[10]: # Sanity check: verify sources are loaded before running expensive NLI
non_empty = sum(1 for s in sources if len(s) > 50)
print(f'Pre-flight check: {non_empty}/{len(sources)} sources have >50_
↳characters')
assert non_empty >= len(sources) * 0.9, f"ERROR: Only {non_empty} sources_
↳loaded. Fix the source loading cell above."

# Quick test on first sample to verify NLI is working
test_score = compute_nli_score(sources[0], summaries[0], nli_tokenizer,
↳nli_model, device)
print(f'Test NLI score on first sample: {test_score:.4f}')
assert test_score > 0.0, "ERROR: NLI scoring returned 0.0 on test sample._
↳Something is wrong."

```

```

print('Pre-flight check PASSED. Starting full NLI scoring...\n')

# Compute NLI scores for all summary-source pairs
# This is the most compute-intensive step
print('Computing NLI-based consistency scores...')
print('This will take approximately 30-60 minutes on a T4 GPU.')
print()

nli_scores = []
for i in tqdm(range(len(sources)), desc='NLI Scoring'):
    score = compute_nli_score(
        sources[i], summaries[i],
        nli_tokenizer, nli_model, device
    )
    nli_scores.append(score)

    # Print progress every 200 samples
    if (i + 1) % 200 == 0:
        print(f'  Processed {i+1}/{len(sources)}, mean NLI score so far: {np.
mean(nli_scores):.4f}')

nli_scores = np.array(nli_scores)
print(f'\nNLI-Score mean: {nli_scores.mean():.4f}, std: {nli_scores.std():.4f}')
print(f'NLI-Score range: [{nli_scores.min():.4f}, {nli_scores.max():.4f}]\n')
print(f'Non-zero NLI scores: {np.sum(nli_scores > 0)}/{len(nli_scores)}\n')

```

Pre-flight check: 1600/1600 sources have >50 characters

Test NLI score on first sample: 0.4241

Pre-flight check PASSED. Starting full NLI scoring..

Computing NLI-based consistency scores...

This will take approximately 30-60 minutes on a T4 GPU.

NLI Scoring: 13%| | 201/1600 [01:09<05:18, 4.40it/s]

Processed 200/1600, mean NLI score so far: 0.8246

NLI Scoring: 25%| | 401/1600 [02:13<07:19, 2.73it/s]

Processed 400/1600, mean NLI score so far: 0.8312

NLI Scoring: 38%| | 601/1600 [03:11<03:37, 4.59it/s]

Processed 600/1600, mean NLI score so far: 0.8301

NLI Scoring: 50%| | 800/1600 [04:15<03:43, 3.58it/s]

Processed 800/1600, mean NLI score so far: 0.8288

NLI Scoring: 63%| | 1001/1600 [05:17<03:14, 3.08it/s]

Processed 1000/1600, mean NLI score so far: 0.8338

NLI Scoring: 75%| | 1200/1600 [06:18<03:12, 2.08it/s]

Processed 1200/1600, mean NLI score so far: 0.8389

NLI Scoring: 88%| | 1401/1600 [07:15<01:12, 2.73it/s]

Processed 1400/1600, mean NLI score so far: 0.8437

NLI Scoring: 100%| | 1600/1600 [08:30<00:00, 3.13it/s]

Processed 1600/1600, mean NLI score so far: 0.8447

NLI-Score mean: 0.8447, std: 0.2335

NLI-Score range: [0.0039, 0.9940]

Non-zero NLI scores: 1600/1600

1.6 6. Compute ConsistEval Composite Score

ConsistEval combines the surface-level BERTScore with the NLI consistency score:

$$\text{ConsistEval}(d, s) = \alpha \cdot f_{\text{surface}}(d, s) + (1 - \alpha) \cdot \text{NLI-Score}(d, s)$$

We use $\alpha = 0.3$ (giving 70% weight to factual consistency).

```
[11]: # Normalise BERTScore to [0, 1] for combination
bertscore_normalised = (bertscore_f1 - bertscore_f1.min()) / (bertscore_f1.
    ↪max() - bertscore_f1.min() + 1e-8)

# Compute ConsistEval with alpha = 0.3
ALPHA = 0.3
consisteval_scores = ALPHA * bertscore_normalised + (1 - ALPHA) * nli_scores

print(f'ConsistEval mean: {consisteval_scores.mean():.4f}, std:␣
    ↪{consisteval_scores.std():.4f}')
print(f'Alpha = {ALPHA} (BERTScore weight = {ALPHA}, NLI weight = {1-ALPHA})')
```

ConsistEval mean: 0.7482, std: 0.1664

Alpha = 0.3 (BERTScore weight = 0.3, NLI weight = 0.7)

1.7 7. Calculate Spearman Correlations with Human Ratings

We now measure how well each automated metric correlates with the human expert consistency ratings.

```
[12]: def bootstrap_spearman(metric_scores, human_scores, n_bootstrap=1000, ci=0.95):
    """
    Compute Spearman correlation with bootstrap confidence intervals.

    Args:
```

```

    metric_scores: Array of automated metric scores
    human_scores: Array of human ratings
    n_bootstrap: Number of bootstrap resamples
    ci: Confidence interval level

Returns:
    tuple: (correlation, lower_bound, upper_bound)
"""
rho, p_value = spearmanr(metric_scores, human_scores)

# Bootstrap confidence intervals
n = len(metric_scores)
bootstrap_rhos = []
rng = np.random.RandomState(42)
for _ in range(n_bootstrap):
    idx = rng.choice(n, size=n, replace=True)
    r, _ = spearmanr(metric_scores[idx], human_scores[idx])
    bootstrap_rhos.append(r)

lower = np.percentile(bootstrap_rhos, (1 - ci) / 2 * 100)
upper = np.percentile(bootstrap_rhos, (1 + ci) / 2 * 100)

return rho, lower, upper, p_value

# Compute correlations for all metrics
metrics = {
    'ROUGE-1': rouge1_scores,
    'ROUGE-2': rouge2_scores,
    'ROUGE-L': rougeL_scores,
    'BERTScore': bertscore_f1,
    'NLI-Score (ours)': nli_scores,
    'ConsistEval (ours)': consisteval_scores,
}

print('=' * 70)
print(f'{"Metric":<25} {"Spearman rho":>12} {"95% CI":>20} {"p-value":>12}')
print('=' * 70)

results_data = []
for name, scores in metrics.items():
    rho, lower, upper, p_val = bootstrap_spearman(scores, human_consistency)
    print(f'{"name":<25} {"rho":>12.4f} ({lower:.4f}, {upper:.4f}) {"p_val":>12.2e}')
    results_data.append({
        'Metric': name,
        'Spearman rho': rho,
        'CI Lower': lower,

```

```

        'CI Upper': upper,
        'p-value': p_val
    })

print('=' * 70)

results_df = pd.DataFrame(results_data)
print('\nRelative improvement of ConsistEval over BERTScore:')
bert_rho = results_df[results_df['Metric'] == 'BERTScore']['Spearman rho'].
    values[0]
ce_rho = results_df[results_df['Metric'] == 'ConsistEval (ours)']['Spearman
    rho'].values[0]
print(f' ({ce_rho:.4f} - {bert_rho:.4f}) / {bert_rho:.4f} = {(ce_rho -
    bert_rho) / abs(bert_rho) * 100:.1f}% improvement')

```

```
=====
```

Metric	Spearman rho	95% CI	p-value
=====			
ROUGE-1	0.1558 (0.1061, 0.2052)		3.68e-10
ROUGE-2	0.1422 (0.0940, 0.1964)		1.13e-08
ROUGE-L	0.1385 (0.0910, 0.1877)		2.65e-08
BERTScore	0.0967 (0.0478, 0.1453)		1.06e-04
NLI-Score (ours)	0.4078 (0.3647, 0.4477)		3.86e-65
ConsistEval (ours)	0.3994 (0.3559, 0.4407)		2.46e-62
=====			

Relative improvement of ConsistEval over BERTScore:
 $(0.3994 - 0.0967) / 0.0967 = 312.9\%$ improvement

1.8 8. Visualisations

1.8.1 8.1 Results Table (for the paper)

```

[13]: # Create a publication-quality results table
print('\n--- Table 1: Spearman correlation with human expert consistency
    ratings ---\n')
for _, row in results_df.iterrows():
    ci_str = f"({row['CI Lower']:.2f}, {row['CI Upper']:.2f})"
    print(f"{row['Metric']:<25} {row['Spearman rho']:.2f} {ci_str}")

```

--- Table 1: Spearman correlation with human expert consistency ratings ---

ROUGE-1	0.16 (0.11, 0.21)
ROUGE-2	0.14 (0.09, 0.20)
ROUGE-L	0.14 (0.09, 0.19)
BERTScore	0.10 (0.05, 0.15)
NLI-Score (ours)	0.41 (0.36, 0.45)

ConsistEval (ours) 0.40 (0.36, 0.44)


1.8.2 8.2 Scatter Plot: Metric Scores vs Human Consistency Ratings

```
[14]: fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Left panel: BERTScore vs Human Consistency
ax1 = axes[0]
ax1.scatter(human_consistency, bertscore_f1, alpha=0.3, s=10, c='steelblue',
            label='BERTScore')
rho_bert = spearmanr(bertscore_f1, human_consistency)[0]
ax1.set_xlabel('Human Consistency Rating', fontsize=12)
ax1.set_ylabel('BERTScore F1', fontsize=12)
ax1.set_title(f'BERTScore vs Human Consistency (  $\rho$  = {rho_bert:.3f})',
             fontsize=13)
ax1.grid(True, alpha=0.3)

# Right panel: ConsistEval vs Human Consistency
ax2 = axes[1]
ax2.scatter(human_consistency, consisteval_scores, alpha=0.3, s=10,
            c='crimson', label='ConsistEval')
rho_ce = spearmanr(consisteval_scores, human_consistency)[0]
ax2.set_xlabel('Human Consistency Rating', fontsize=12)
ax2.set_ylabel('ConsistEval Score', fontsize=12)
ax2.set_title(f'ConsistEval vs Human Consistency (  $\rho$  = {rho_ce:.3f})',
             fontsize=13)
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig('scatter_plot.png', dpi=300, bbox_inches='tight')
plt.show()
print('Scatter plot saved to scatter_plot.png')
```



Scatter plot saved to scatter_plot.png

1.8.3 8.3 Bar Chart: Correlation Comparison Across All Metrics

```
[15]: fig, ax = plt.subplots(figsize=(10, 5))


metric_names = results_df['Metric'].tolist()
rhos = results_df['Spearman rho'].tolist()
ci_lower = results_df['CI Lower'].tolist()
ci_upper = results_df['CI Upper'].tolist()

# Error bars
errors = [[r - l for r, l in zip(rhos, ci_lower)],
          [u - r for r, u in zip(rhos, ci_upper)]]

colours = ['#4e79a7', '#4e79a7', '#4e79a7', '#4e79a7', '#e15759', '#e15759']
bars = ax.bar(metric_names, rhos, color=colours, edgecolor='black', linewidth=0.
↳5,
              yerr=errors, capsize=4, error_kw={'linewidth': 1})

ax.set_ylabel('Spearman Correlation ( )', fontsize=12)
ax.set_title('Correlation with Human Expert Consistency Ratings', fontsize=13)
ax.set_ylim(0, 1.0)
ax.axhline(y=0.4, color='grey', linestyle='--', alpha=0.5, label=' = 0.4↳
↳threshold')
ax.legend(fontsize=10)
plt.xticks(rotation=30, ha='right', fontsize=10)
ax.grid(axis='y', alpha=0.3)

plt.tight_layout()
plt.savefig('bar_chart.png', dpi=300, bbox_inches='tight')
plt.show()
print('Bar chart saved to bar_chart.png')
```



Bar chart saved to bar_chart.png

1.8.4 8.4 Correlation with All Human Dimensions

For additional analysis, we also check how the metrics correlate with other human evaluation dimensions (coherence, fluency, relevance) to confirm that our metric specifically targets consistency.

```
[16]: # Cross-dimension correlation analysis
dimensions = {
    'Consistency': human_consistency,
    'Coherence': human_coherence,
    'Fluency': human_fluency,
    'Relevance': human_relevance,
}

cross_results = []
for metric_name, metric_scores in metrics.items():
    row = {'Metric': metric_name}
    for dim_name, dim_scores in dimensions.items():
        rho, _ = spearmanr(metric_scores, dim_scores)
        row[dim_name] = rho
    cross_results.append(row)

cross_df = pd.DataFrame(cross_results)
cross_df.set_index('Metric', inplace=True)

print('\n--- Cross-Dimension Correlation Analysis ---')
print('Spearman correlation between each metric and each human evaluation_
↪dimension:\n')
```

```

print(cross_df.round(3).to_string())


# Heatmap
fig, ax = plt.subplots(figsize=(8, 5))
sns.heatmap(cross_df.astype(float), annot=True, fmt='.3f', cmap='RdYlGn',
            vmin=-0.1, vmax=0.8, ax=ax, linewidths=0.5)
ax.set_title('Metric-Human Dimension Correlation Matrix', fontsize=13)
plt.tight_layout()
plt.savefig('heatmap.png', dpi=300, bbox_inches='tight')
plt.show()
print('Heatmap saved to heatmap.png')

```

--- Cross-Dimension Correlation Analysis ---

Spearman correlation between each metric and each human evaluation dimension:

	Consistency	Coherence	Fluency	Relevance
Metric				
ROUGE-1	0.156	0.172	0.096	0.331
ROUGE-2	0.142	0.134	0.085	0.257
ROUGE-L	0.139	0.169	0.107	0.268
BERTScore	0.097	0.329	0.126	0.356
NLI-Score (ours)	0.408	0.168	0.344	0.197
ConsistEval (ours)	0.399	0.237	0.328	0.290



Heatmap saved to heatmap.png

1.9 9. Error Analysis

We examine cases where ConsistEval disagrees most with human ratings to understand its strengths and failure modes.

```
[17]: # Normalise ConsistEval to 1-5 scale for comparison
ce_normalised = 1 + 4 * (consisteval_scores - consisteval_scores.min()) / \
    (consisteval_scores.max() - consisteval_scores.min() + 1e-8)

# Compute disagreement
disagreement = np.abs(ce_normalised - human_consistency)

# Find worst disagreements
worst_idx = np.argsort(disagreement)[::-1][:20]

print('Top 20 cases where ConsistEval disagrees most with human ratings:\n')
print(f'{"Idx":>5} {"Human":>7} {"CE":>7} {"Gap":>7} {"NLI":>7} {"BERT":>7} ␣
    ↳Summary (first 80 chars)')
print('-' * 120)

for idx in worst_idx:
    summ_preview = summaries[idx][:80].replace('\n', ' ')
    print(f'{"idx":>5} {"human_consistency[idx]:>7.2f} {"ce_normalised[idx]:>7.2f} '
        f'{"disagreement[idx]:>7.2f} {"nli_scores[idx]:>7.3f}␣
    ↳{"bertscore_f1[idx]:>7.3f} {"summ_preview}')
```

Top 20 cases where ConsistEval disagrees most with human ratings:

Idx	Human	CE	Gap	NLI	BERT	Summary (first 80 chars)
1048	5.00	1.26	3.74	0.007	0.855	Dyke wants to increase the number of home-grown players at top flight clubs from
1540	5.00	1.38	3.62	0.065	0.848	serge gnabry says he is feeling better and better ' all the time as the arsena
728	5.00	1.50	3.50	0.066	0.862	Chipotle is planning to launch a delivery service for online and mobile orders i
467	5.00	1.51	3.49	0.031	0.876	new zealanders tim weston was on holiday in puerto vallarta with his wife when th
1240	5.00	1.52	3.48	0.071	0.862	Dayna Dobias, 19, from Downers Grove, Illinois, is attempting to raise awareness
744	5.00	1.56	3.44	0.078	0.864	The Boston Marathon bombing verdict is a resounding victory for the victims' fam
1320	5.00	1.57	3.43	0.071	0.868	An FBI agent's sniper rifle was stolen from his car during a hotel parking lot r
631	5.00	1.57	3.43	0.100	0.858	nathan hughes's knee collided

with george north's head as he scored a try for no						
1560	5.00	1.58	3.42	0.015	0.890	Jacob Phillips fell down a 70ft cliff to his death in Cardiff.
1191	5.00	1.61	3.39	0.234	0.815	arsenal host chelsea at the emirates stadium on sunday (4pm kick-off) mesut ozil
936	5.00	1.63	3.37	0.179	0.836	Vanessa Santillan worked as a transgender escort in London. She was found strang
131	5.00	1.63	3.37	0.045	0.885	visitors at cleveland metroparks zoo in ohio heard a scream as the toddler tumbl
138	5.00	1.63	3.37	0.045	0.885	visitors at cleveland metroparks zoo in ohio heard a scream as the toddler tumbl
134	5.00	1.63	3.37	0.045	0.885	visitors at cleveland metroparks zoo in ohio heard a scream as the toddler tumbl
130	5.00	1.63	3.37	0.045	0.885	visitors at cleveland metroparks zoo in ohio heard a scream as the toddler tumbl
184	5.00	1.69	3.31	0.124	0.863	Barcelona team-mates Neymar and Dani Alves head to watch El Clasico on Thursday
129	5.00	1.69	3.31	0.061	0.886	cleveland metroparks zoo in ohio heard a scream as the toddler tumbled into the
1256	5.00	1.71	3.29	0.095	0.875	Javier Hernandez is in the final year of his contract at Real Madrid and is look
1351	5.00	1.72	3.28	0.112	0.870	everton beat burnley 1-0 at goodison park on saturday . ross barkley missed from
793	5.00	1.72	3.28	0.020	0.903	the 69-year-old actress collaborated with nbc's today show to launch a contest f

```
[18]: # Detailed inspection of specific disagreement cases
print('\n=== Detailed Error Analysis: Top 5 Disagreement Cases ===\n')

for rank, idx in enumerate(worst_idx[:5], 1):
    print(f'--- Case {rank} (Index {idx}) ---')
    print(f'Human Consistency: {human_consistency[idx]:.2f}')
    print(f'ConsistEval Score: {consisteval_scores[idx]:.4f} (normalised: {ce_normalised[idx]:.2f})')
    print(f'NLI-Score: {nli_scores[idx]:.4f}')
    print(f'BERTScore: {bertscore_f1[idx]:.4f}')
    print(f'ROUGE-1: {rouge1_scores[idx]:.4f}')
    print(f'Model: {model_ids[idx]}')
    print(f'\nSource (first 300 chars):\n{sources[idx][:300]}...')
    print(f'\nSummary:\n{summaries[idx][:500]}')
    print(f'\n{"="*80}\n')
```

=== Detailed Error Analysis: Top 5 Disagreement Cases ===

```

--- Case 1 (Index 1048) ---
Human Consistency: 5.00
ConsistEval Score: 0.1442 (normalised: 1.26)

```

NLI-Score: 0.0068
BERTScore: 0.8554
ROUGE-1: 0.3088
Model: M20

Source (first 300 chars):

Football Association chairman Greg Dyke has been handed the backing of five former England managers as he bids to push through radical changes. Dyke wants to increase the minimum number of home-grown players in club squads from eight to 12, however he is facing opposition from the Premier League. Th...

Summary:

Dyke wants to increase the number of home-grown players at top flight clubs from eight to 12 'The current crop of young players is the most talented ever assembled by the game and has the potential to be the best in England's history,' they write. 'We are all indebted to the England coaches who have helped us achieve this. 'The current crop of young players is the most talented ever assembled by the game and has the potential to be the best in England

=====

--- Case 2 (Index 1540) ---

Human Consistency: 5.00
ConsistEval Score: 0.1708 (normalised: 1.38)
NLI-Score: 0.0646
BERTScore: 0.8484
ROUGE-1: 0.4356
Model: M15

Source (first 300 chars):

Serge Gnabry says he is feeling 'better and better' all the time as the Arsenal teenager prepares to return after more than a year out. The 19-year-old midfielder impressed for the Gunners last season, but his progress has been stunted by a serious knee injury. Gnabry, as a result, has not featured ...

Summary:

serge gnabry says he is feeling ` better and better ' all the time as the arsenal teenager impressed for the gunners last season . the 19-year-old has not featured since the champions league last-16 second leg against bayern munich in march 2014 . the germany under-21 international is at last closing in on a return .

=====

--- Case 3 (Index 728) ---

Human Consistency: 5.00
ConsistEval Score: 0.1986 (normalised: 1.50)
NLI-Score: 0.0664

BERTScore: 0.8619
ROUGE-1: 0.2632
Model: M20

Source (first 300 chars):

Mexican restaurant Chipotle has decided to tap into the \$70 billion food delivery market by teaming up with an app to bring burritos straight to customers' doors. The fast-casual chain will work with the Postmates app to begin offering delivery for online and mobile orders in 67 cities, according to...

Summary:

Chipotle is planning to launch a delivery service for online and mobile orders in 67 cities by the end of the year

=====

--- Case 4 (Index 467) ---

Human Consistency: 5.00
ConsistEval Score: 0.2006 (normalised: 1.51)
NLI-Score: 0.0307
BERTScore: 0.8756
ROUGE-1: 0.4800
Model: M14

Source (first 300 chars):

A video of a menacing crocodile swimming through a public marina with a dead pet dog locked in between its jaws has been viewed more than half a million times online. The clip shows a large croc calmly moving through the marina in the tourist hotspot of Puerto Vallarta, in Mexico, with a black dog w...

Summary:

new zealander tim weston was on holiday in puerto vallarta with his wife when the attack happened and described how ` the crocodile kept the dog in its mouth for ages , not moving at all ' he told daily mail australia the dog belonged to the local bike shop 's owner .

=====

--- Case 5 (Index 1240) ---

Human Consistency: 5.00
ConsistEval Score: 0.2028 (normalised: 1.52)
NLI-Score: 0.0712
BERTScore: 0.8623
ROUGE-1: 0.3030
Model: M20

Source (first 300 chars):

A teenager from Illinois is tackling her disability head-on and attempting to positively influence thousands of others by dancing. Dayna Dobias, 19, from Downers Grove was born with cerebral palsy, but she hopes to inspire others with her enthusiastic videos. 'I've gotten bullied because of it and I...

Summary:

Dayna Dobias, 19, from Downers Grove, Illinois, is attempting to raise awareness of her disability VIDEO: Dayna Dobias, 18, from Downers Grove, Illinois, is trying to raise awareness of her disability She says she has been bullied at school and has been denied a job because of her disability. 'I've been bullied for my disability and I've been kicked out of school and I've been told I'm not really disabled,' she said

=====

```
[19]: # Categorise disagreement types
# Cases where ConsistEval is LOW but human says HIGH (false negatives)
false_negatives = np.where((ce_normalised < 3) & (human_consistency >= 4))[0]
# Cases where ConsistEval is HIGH but human says LOW (false positives)
false_positives = np.where((ce_normalised >= 4) & (human_consistency < 3))[0]
# Cases where both agree on HIGH
true_positives = np.where((ce_normalised >= 4) & (human_consistency >= 4))[0]
# Cases where both agree on LOW
true_negatives = np.where((ce_normalised < 3) & (human_consistency < 3))[0]

print('=== ConsistEval Agreement Analysis ===')
print(f'True Positives (both high): {len(true_positives):>5}┐
      ↳ ({len(true_positives)/len(sources)*100:.1f}%)')
print(f'True Negatives (both low): {len(true_negatives):>5}┐
      ↳ ({len(true_negatives)/len(sources)*100:.1f}%)')
print(f'False Negatives (CE low, H high): {len(false_negatives):>5}┐
      ↳ ({len(false_negatives)/len(sources)*100:.1f}%)')
print(f'False Positives (CE high, H low): {len(false_positives):>5}┐
      ↳ ({len(false_positives)/len(sources)*100:.1f}%)')
print(f'\nOverall agreement rate: {(len(true_positives) + len(true_negatives)) /
      ↳ len(sources) * 100:.1f}%)')
```

```
=== ConsistEval Agreement Analysis ===
True Positives (both high): 1028 (64.2%)
True Negatives (both low): 79 (4.9%)
False Negatives (CE low, H high): 107 (6.7%)
False Positives (CE high, H low): 16 (1.0%)
```

Overall agreement rate: 69.2%

1.10 10. Alpha Sensitivity Analysis

We test different values of α to understand how the balance between BERTScore and NLI-Score affects the overall correlation.


```
[20]: alphas = np.arange(0.0, 1.05, 0.05)
alpha_rhos = []

for a in alphas:
    composite = a * bertscore_normalised + (1 - a) * nli_scores
    rho, _ = spearmanr(composite, human_consistency)
    alpha_rhos.append(rho)

best_alpha = alphas[np.argmax(alpha_rhos)]
best_rho = max(alpha_rhos)

fig, ax = plt.subplots(figsize=(8, 4))
ax.plot(alphas, alpha_rhos, 'b-o', markersize=4, linewidth=1.5)
ax.axvline(x=ALPHA, color='red', linestyle='--', alpha=0.7, label=f'Selected  $\alpha$  = {ALPHA}')
ax.axvline(x=best_alpha, color='green', linestyle='--', alpha=0.7, label=f'Best  $\alpha$  = {best_alpha:.2f} (  $\rho$  = {best_rho:.4f})')
ax.set_xlabel(' (BERTScore weight)', fontsize=12)
ax.set_ylabel('Spearman  $\rho$  with Human Consistency', fontsize=12)
ax.set_title('Sensitivity Analysis: Effect of  $\alpha$  on ConsistEval Correlation',
             ↪fontsize=13)
ax.legend(fontsize=10)
ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig('alpha_sensitivity.png', dpi=300, bbox_inches='tight')
plt.show()
print(f'\nBest alpha: {best_alpha:.2f} with Spearman rho = {best_rho:.4f}')
print(f'Selected alpha: {ALPHA} with Spearman rho = {alpha_rhos[int(ALPHA/0.05)]:.4f}')
```



Best alpha: 0.00 with Spearman rho = 0.4078
 Selected alpha: 0.3 with Spearman rho = 0.4016

1.11 11. Per-Model Analysis

The tutor feedback requires showing metric scores for **every model** in SummEval (all 16 summarisation systems). We group them into three architectural families: - **Statistical Extractive** (6 models): Lead-3, TextRank, LexRank, Luhn, SumBasic, KLSum - **Neural Seq2Seq** (5 models): Pointer-Gen, PGN+Coverage, Fast-Abs-RL, Bottom-Up, Improve-Abs - **Transformer Pre-trained** (5 models): Unified-Ext-Abs, ROOGe, SENECA, BertSumExt, BertSumExtAbs, BART, Pegasus, T5

This allows us to investigate whether ROUGE vs NLI-based metrics behave differently across model architectures.

```
[21]: # Discover the unique model IDs present in the SummEval data
unique_models = sorted(set(model_ids))
print(f'Unique model IDs ({len(unique_models)}): {unique_models}')

# Count samples per model
from collections import Counter
model_counts = Counter(model_ids)
for mid in unique_models:
    print(f' {mid}: {model_counts[mid]} samples')
```

```
Unique model IDs (16): ['M0', 'M1', 'M10', 'M11', 'M12', 'M13', 'M14', 'M15',
'M17', 'M2', 'M20', 'M22', 'M23', 'M5', 'M8', 'M9']
M0: 100 samples
M1: 100 samples
```

M10: 100 samples
 M11: 100 samples
 M12: 100 samples
 M13: 100 samples
 M14: 100 samples
 M15: 100 samples
 M17: 100 samples
 M2: 100 samples
 M20: 100 samples
 M22: 100 samples
 M23: 100 samples
 M5: 100 samples
 M8: 100 samples
 M9: 100 samples

```

[22]: # Map SummEval model IDs to human-readable names
# Reference: Fabbri et al. (2021) Table 1 and supplementary materials
MODEL_NAMES = {
    'M0': 'Lead-3',
    'M1': 'TextRank',
    'M2': 'LexRank',
    'M5': 'SumBasic',
    'M8': 'Pointer-Gen',
    'M9': 'PGN+Coverage',
    'M10': 'Fast-Abs-RL',
    'M11': 'Bottom-Up',
    'M12': 'Improve-Abs',
    'M13': 'Unified-Ext-Abs',
    'M14': 'ROOGe',
    'M15': 'SENECA',
    'M17': 'BertSumExt',
    'M20': 'BART',
    'M22': 'Pegasus',
    'M23': 'T5',
}

# Architectural family categories
STATISTICAL = {'M0', 'M1', 'M2', 'M5'} # Lead-3, TextRank, LexRank, SumBasic
NEURAL_SEQ2SEQ = {'M8', 'M9', 'M10', 'M11', 'M12'} # Pointer-Gen through Improve-Abs
TRANSFORMER = {'M13', 'M14', 'M15', 'M17', 'M20', 'M22', 'M23'} # Pretrained transformer models

def get_category(model_id):
    if model_id in STATISTICAL:
        return 'Statistical Extractive'

```

```

elif model_id in NEURAL_SEQ2SEQ:
    return 'Neural Seq2Seq'
elif model_id in TRANSFORMER:
    return 'Transformer Pretrained'
else:
    return 'Unknown'

def get_name(model_id):
    return MODEL_NAMES.get(model_id, model_id)

# Verify all model IDs are mapped
unmapped = [m for m in unique_models if m not in MODEL_NAMES]
if unmapped:
    print(f'WARNING: Unmapped model IDs: {unmapped}')
else:
    print(f'All {len(unique_models)} model IDs successfully mapped.')

# Show the mapping
for mid in sorted(MODEL_NAMES.keys(), key=lambda x: int(x[1:])):
    print(f' {mid} -> {get_name(mid)} [{get_category(mid)}]')

```

All 16 model IDs successfully mapped.

```

M0 -> Lead-3 [Statistical Extractive]
M1 -> TextRank [Statistical Extractive]
M2 -> LexRank [Statistical Extractive]
M5 -> SumBasic [Statistical Extractive]
M8 -> Pointer-Gen [Neural Seq2Seq]
M9 -> PGN+Coverage [Neural Seq2Seq]
M10 -> Fast-Abs-RL [Neural Seq2Seq]
M11 -> Bottom-Up [Neural Seq2Seq]
M12 -> Improve-Abs [Neural Seq2Seq]
M13 -> Unified-Ext-Abs [Transformer Pretrained]
M14 -> ROOGe [Transformer Pretrained]
M15 -> SENECA [Transformer Pretrained]
M17 -> BertSumExt [Transformer Pretrained]
M20 -> BART [Transformer Pretrained]
M22 -> Pegasus [Transformer Pretrained]
M23 -> T5 [Transformer Pretrained]

```

```

[23]: # Build a DataFrame with per-sample metrics and model info
sample_df = pd.DataFrame({
    'model_id': model_ids,
    'model_name': [get_name(m) for m in model_ids],
    'category': [get_category(m) for m in model_ids],
    'human_consistency': human_consistency,
    'rouge1': rouge1_scores,
    'rouge2': rouge2_scores,

```

```

        'rougeL': rougeL_scores,
        'bertscore': bertscore_f1,
        'nli_score': nli_scores,
        'consisteval': consisteval_scores,
    })

# Group by model and compute means
per_model_df = sample_df.groupby(['model_id', 'model_name', 'category']).agg({
    'human_consistency': 'mean',
    'rouge1': 'mean',
    'rouge2': 'mean',
    'rougeL': 'mean',
    'bertscore': 'mean',
    'nli_score': 'mean',
    'consisteval': 'mean',
}).reset_index()

# Sort by category then model_id for display
category_order = {'Statistical Extractive': 0, 'Neural Seq2Seq': 1,
    ↪ 'Transformer Pretrained': 2}
per_model_df['cat_order'] = per_model_df['category'].map(category_order)
per_model_df = per_model_df.sort_values(['cat_order', 'model_id']).
    ↪ drop(columns='cat_order')

# Display the table (this is what goes into Table 2 of the paper)
print('\n--- Table 2: Per-Model Metric Scores (for paper) ---\n')
print(f'{"Model":<22} {"Cat.":<6} {"Cons.":>5} {"R-1":>6} {"R-2":>6} {"R-L":>6}
    ↪ {"BERT":>6} {"NLI":>6} {"CE":>6}')
print('-' * 75)
for _, row in per_model_df.iterrows():
    cat_abbr = {'Statistical Extractive': 'Stat', 'Neural Seq2Seq': 'Neur',
    ↪ 'Transformer Pretrained': 'Trans'}
    print(f"{row['model_name']:<22} {cat_abbr[row['category']]:<6} "
        f"{row['human_consistency']:5.2f} "
        f"{row['rouge1']:6.4f} {row['rouge2']:6.4f} {row['rougeL']:6.4f} "
        f"{row['bertscore']:6.4f} {row['nli_score']:6.4f} {row['consisteval']:
    ↪ 6.4f}")

# Category-level means
print('\n--- Category Means ---\n')
cat_means = per_model_df.
    ↪ groupby('category')[['human_consistency', 'rouge1', 'rouge2', 'rougeL', 'bertscore', 'nli_score']
    ↪ mean()
for cat in ['Statistical Extractive', 'Neural Seq2Seq', 'Transformer
    ↪ Pretrained']:
    r = cat_means.loc[cat]

```

```
print(f"{cat:<25} Cons={r['human_consistency']:.2f} R-1={r['rouge1']:.4f} NLI={r['nli_score']:.4f} CE={r['consisteval']:.4f}")
```

--- Table 2: Per-Model Metric Scores (for paper) ---

Model	Cat.	Cons.	R-1	R-2	R-L	BERT	NLI	CE
Lead-3	Stat	4.98	0.4187	0.1927	0.2678	0.8597	0.9692	0.8262
TextRank	Stat	4.98	0.4245	0.1988	0.2731	0.8608	0.9641	0.8248
LexRank	Stat	4.99	0.4212	0.1908	0.2700	0.8597	0.9660	0.8239
SumBasic	Stat	4.97	0.4306	0.2024	0.2746	0.8613	0.9358	0.8061
Fast-Abs-RL	Neur	4.25	0.4447	0.2133	0.3028	0.8616	0.7590	0.6828
Bottom-Up	Neur	3.27	0.4139	0.1835	0.2688	0.8562	0.6708	0.6104
Improve-Abs	Neur	4.96	0.4246	0.1959	0.2828	0.8610	0.9514	0.8163
Pointer-Gen	Neur	4.65	0.4209	0.1955	0.2903	0.8624	0.9032	0.7853
PGN+Coverage	Neur	4.67	0.4281	0.1951	0.2743	0.8597	0.9377	0.8041
Unified-Ext-Abs	Trans	4.82	0.4247	0.1976	0.2958	0.8606	0.9013	0.7805
ROOGe	Trans	4.90	0.4192	0.1996	0.2939	0.8615	0.8945	0.7775
SENECA	Trans	4.94	0.4206	0.1942	0.2857	0.8597	0.8889	0.7700
BertSumExt	Trans	4.93	0.4753	0.2343	0.3301	0.8734	0.7442	0.6955
BART	Trans	3.40	0.3040	0.1032	0.2107	0.8716	0.4240	0.4680
Pegasus	Trans	4.94	0.4594	0.2287	0.3284	0.8877	0.8156	0.7736
T5	Trans	4.91	0.4755	0.2418	0.3413	0.8732	0.7895	0.7269

--- Category Means ---

Statistical Extractive	Cons=4.98	R-1=0.4238	NLI=0.9588	CE=0.8202
Neural Seq2Seq	Cons=4.36	R-1=0.4265	NLI=0.8444	CE=0.7398
Transformer Pretrained	Cons=4.69	R-1=0.4255	NLI=0.7797	CE=0.7132

```
[24]: # Visualisation: Per-model metric scores (3-panel horizontal bar chart)
fig, axes = plt.subplots(1, 3, figsize=(18, 8), sharey=True)

# Colour-code by category
colors = {'Statistical Extractive': '#4e79a7', 'Neural Seq2Seq': '#f28e2b',
        'Transformer Pretrained': '#59a14f'}
bar_colors = [colors[cat] for cat in per_model_df['category']]
model_labels = per_model_df['model_name'].tolist()

# Panel 1: ROUGE-1
axes[0].barh(model_labels, per_model_df['rouge1'], color=bar_colors)
axes[0].set_xlabel('ROUGE-1 F1')
axes[0].set_title('Surface-Level: ROUGE-1')
axes[0].invert_yaxis()

# Panel 2: NLI-Score
```

```


axes[1].barh(model_labels, per_model_df['nli_score'], color=bar_colors)
axes[1].set_xlabel('NLI-Score')
axes[1].set_title('Semantic: NLI-Score')

# Panel 3: Human Consistency
axes[2].barh(model_labels, per_model_df['human_consistency'], color=bar_colors)
axes[2].set_xlabel('Human Consistency (1-5)')
axes[2].set_title('Human Expert Rating')

# Add legend
from matplotlib.patches import Patch
legend_elements = [Patch(facecolor=c, label=l) for l, c in colors.items()]
fig.legend(handles=legend_elements, loc='lower center', ncol=3, fontsize=11,
           bbox_to_anchor=(0.5, -0.02))

plt.suptitle('Per-Model Metric Scores Across Architectural Families',
             fontsize=14, fontweight='bold')
plt.tight_layout(rect=[0, 0.03, 1, 0.96])
plt.savefig('per_model_scores.png', dpi=150, bbox_inches='tight')
plt.show()
print('Saved: per_model_scores.png')

```



Saved: per_model_scores.png

1.11.1 11.1 Architecture-Level Correlation Analysis

We compute Spearman correlations **within each architectural family** to test the hypothesis that ROUGE correlates better with statistical models while NLI-based metrics correlate better with transformer-based models.

```
[25]: # Compute within-group Spearman correlations
from scipy.stats import spearmanr

print('\n--- Within-Group Spearman Correlations (metric vs human consistency)␣
↳---\n')
print(f'{"Category":<27} {"ROUGE-1":>10} {"NLI-Score":>10} {"ConsistEval":>12}')
print('-' * 62)

for cat in ['Statistical Extractive', 'Neural Seq2Seq', 'Transformer␣
↳Pretrained']:
    mask = sample_df['category'] == cat
    subset = sample_df[mask]
    n = len(subset)

    r1_rho, r1_p = spearmanr(subset['rouge1'], subset['human_consistency'])
    nli_rho, nli_p = spearmanr(subset['nli_score'], subset['human_consistency'])
    ce_rho, ce_p = spearmanr(subset['consisteval'], subset['human_consistency'])

    def fmt(rho, p):
        sig = '***' if p < 0.001 else '**' if p < 0.01 else '*' if p < 0.05␣
↳else ''
        return f'{rho:6.3f}-{sig}'

    print(f'{cat:<27} {fmt(r1_rho, r1_p):>10} {fmt(nli_rho, nli_p):>10}␣
↳{fmt(ce_rho, ce_p):>12} (n={n})')

print('\nSignificance: *** p<0.001, ** p<0.01, * p<0.05')
```

--- Within-Group Spearman Correlations (metric vs human consistency) ---

Category	ROUGE-1	NLI-Score	ConsistEval	
Statistical Extractive	0.086	0.178***	0.120*	(n=400)
Neural Seq2Seq	0.108*	0.531***	0.509***	(n=500)
Transformer Pretrained	0.251***	0.375***	0.381***	(n=700)

Significance: *** p<0.001, ** p<0.01, * p<0.05

```
[26]: # Print key observations for Section 6.3 of the paper
print('\n=== KEY OBSERVATIONS FOR PAPER (Section 6.3: Architecture-Level␣
↳Analysis) ===\n')

# Category means
for cat in ['Statistical Extractive', 'Neural Seq2Seq', 'Transformer␣
↳Pretrained']:
    subset = sample_df[sample_df['category'] == cat]
```



```

print(f'{cat}:')
print(f'  Mean Human Consistency: {subset["human_consistency"].mean():.2f}')
print(f'  Mean ROUGE-1:           {subset["rouge1"].mean():.4f}')
print(f'  Mean NLI-Score:          {subset["nli_score"].mean():.4f}')
print(f'  Mean ConsistEval:        {subset["consisteval"].mean():.4f}')
print()

# ROUGE bias analysis
stat_r1 = sample_df[sample_df['category'] == 'Statistical_
↳Extractive']['rouge1'].mean()
trans_r1 = sample_df[sample_df['category'] == 'Transformer_
↳Pretrained']['rouge1'].mean()
stat_cons = sample_df[sample_df['category'] == 'Statistical_
↳Extractive']['human_consistency'].mean()
trans_cons = sample_df[sample_df['category'] == 'Transformer_
↳Pretrained']['human_consistency'].mean()

print(f'ROUGE-1 gap (Trans - Stat): {trans_r1 - stat_r1:+.4f}')
print(f'Human Consistency gap (Trans - Stat): {trans_cons - stat_cons:+.2f}')
print(f'=> ROUGE overestimates transformer advantage by {abs((trans_r1 -
↳stat_r1) / stat_r1 * 100):.1f}%',)
print(f'  while humans rate them similarly ({stat_cons:.2f} vs {trans_cons:.
↳2f}%',)

# NLI tracking analysis
stat_nli = sample_df[sample_df['category'] == 'Statistical_
↳Extractive']['nli_score'].mean()
trans_nli = sample_df[sample_df['category'] == 'Transformer_
↳Pretrained']['nli_score'].mean()
print(f'\nNLI-Score gap (Trans - Stat): {trans_nli - stat_nli:+.4f}')
print(f'=> NLI tracks human consistency more faithfully across architectures')

print('\n=== UPDATE TABLE 2 IN main.tex WITH THE PER-MODEL VALUES ABOVE ===')

```

=== KEY OBSERVATIONS FOR PAPER (Section 6.3: Architecture-Level Analysis) ===

Statistical Extractive:

Mean Human Consistency:	4.98
Mean ROUGE-1:	0.4238
Mean NLI-Score:	0.9588
Mean ConsistEval:	0.8202

Neural Seq2Seq:

Mean Human Consistency:	4.36
Mean ROUGE-1:	0.4265
Mean NLI-Score:	0.8444

Mean ConsistEval: 0.7398

Transformer Pretrained:

Mean Human Consistency: 4.69
Mean ROUGE-1: 0.4255
Mean NLI-Score: 0.7797
Mean ConsistEval: 0.7132

ROUGE-1 gap (Trans - Stat): +0.0018

Human Consistency gap (Trans - Stat): -0.29

=> ROUGE overestimates transformer advantage by 0.4%
while humans rate them similarly (4.98 vs 4.69)

NLI-Score gap (Trans - Stat): -0.1791

=> NLI tracks human consistency more faithfully across architectures

=== UPDATE TABLE 2 IN main.tex WITH THE PER-MODEL VALUES ABOVE ===

1.12 12. Summary of Findings

Print the final summary that can be used to update the paper with real numbers.

```
[27]: print('=' * 70)
print('FINAL RESULTS SUMMARY FOR PAPER')
print('=' * 70)
print()
print('Table 1 values (update main.tex with these):')
print()
for _, row in results_df.iterrows():
    print(f" {row['Metric']:<25} rho = {row['Spearman rho']:.2f} "
          f"CI = ({row['CI Lower']:.2f}, {row['CI Upper']:.2f})")

print()
print(f'Relative improvement over BERTScore: '
      f'{(ce_rho - bert_rho) / abs(bert_rho) * 100:.1f}%')
print(f'Best alpha from sensitivity analysis: {best_alpha:.2f}')
print(f'Total samples evaluated: {len(sources)}')
print(f'Unique models in SummEval: {len(set(model_ids))}')
print()
print('--- Per-Model Table 2 Data ---')
print(per_model_df[['model_name', 'category', 'human_consistency', 'rouge1', 'rouge2', 'rougeL', 'bert_score', 'nli_score', 'consisteval']].to_string(index=False))
print()
print('Figures saved: scatter_plot.png, bar_chart.png, heatmap.png, '
      'alpha_sensitivity.png, per_model_scores.png')
print('=' * 70)
```

=====

FINAL RESULTS SUMMARY FOR PAPER

=====

Table 1 values (update main.tex with these):

ROUGE-1	rho = 0.16	CI = (0.11, 0.21)
ROUGE-2	rho = 0.14	CI = (0.09, 0.20)
ROUGE-L	rho = 0.14	CI = (0.09, 0.19)
BERTScore	rho = 0.10	CI = (0.05, 0.15)
NLI-Score (ours)	rho = 0.41	CI = (0.36, 0.45)
ConsistEval (ours)	rho = 0.40	CI = (0.36, 0.44)

Relative improvement over BERTScore: 294.2%

Best alpha from sensitivity analysis: 0.00

Total samples evaluated: 1600

Unique models in SummEval: 16

--- Per-Model Table 2 Data ---

model_name	category	human_consistency	rouge1	rouge2
rougeL bertscore nli_score consisteval				
Lead-3 Statistical Extractive		4.983333	0.418675	0.192741
0.267761 0.859658 0.969233 0.826161				
TextRank Statistical Extractive		4.980000	0.424542	0.198776
0.273053 0.860792 0.964079 0.824779				
LexRank Statistical Extractive		4.990000	0.421236	0.190837
0.269972 0.859680 0.965993 0.823935				
SumBasic Statistical Extractive		4.973333	0.430569	0.202393
0.274579 0.861331 0.935848 0.806078				
Fast-Abs-RL Neural Seq2Seq		4.253333	0.444750	0.213340
0.302830 0.861580 0.758972 0.682754				
Bottom-Up Neural Seq2Seq		3.270000	0.413949	0.183527
0.268832 0.856168 0.670849 0.610437				
Improve-Abs Neural Seq2Seq		4.960000	0.424592	0.195935
0.282783 0.860989 0.951391 0.816285				
Pointer-Gen Neural Seq2Seq		4.653333	0.420939	0.195507
0.290325 0.862357 0.903238 0.785266				
PGN+Coverage Neural Seq2Seq		4.673333	0.428065	0.195138
0.274302 0.859675 0.937715 0.804131				
Unified-Ext-Abs Transformer Pretrained		4.820000	0.424679	0.197572
0.295793 0.860629 0.901266 0.780490				
ROOGe Transformer Pretrained		4.896667	0.419218	0.199639
0.293892 0.861528 0.894485 0.777509				
SENECA Transformer Pretrained		4.936667	0.420587	0.194152
0.285652 0.859693 0.888918 0.770009				
BertSumExt Transformer Pretrained		4.926667	0.475316	0.234311
0.330114 0.873356 0.744195 0.695541				
BART Transformer Pretrained		3.396667	0.304000	0.103183

0.210677	0.871634	0.423999	0.468021	
	Pegasus Transformer Pretrained			4.943333 0.459379 0.228668
0.328402	0.887659	0.815636	0.773646	
	T5 Transformer Pretrained			4.910000 0.475489 0.241784
0.341296	0.873163	0.789537	0.726901	

Figures saved: scatter_plot.png, bar_chart.png, heatmap.png,
alpha_sensitivity.png, per_model_scores.png

=====